

CAP5638 Project 2

Classification Using Linear Discriminant Functions and Boosting Algorithms

Suhib Sam Kiswani

December 2, 2015

The algorithms were implemented in *Python 3.5*, with a dependence on the *scipy* [1] library.

1 Basic Two-Class Classification Using Perceptron Algorithms

Abstractly, the problem is as follows: Given n labeled training samples, $D = \{(x_1, L_1), (x_2, L_2), \dots, (x_n, L_n)\}$, where $L_i = \pm 1$, implement Algorithm 4 (Fixed-Increment Single-Sample Perceptron Algorithm) and Algorithm 8 (Batch Relaxation with Margin) of Chapter 5 in the textbook.

The algorithms used for this method are:

Algorithm 5.4 (Fixed-Increment Single-Sample Perceptron)

```
1: initialize  $a, k = 0$ 
2: do  $k \leftarrow (k + 1) \bmod n$ 
3:   if  $\mathbf{y}_k$  is misclassified by  $\mathbf{a}$  then  $\mathbf{a} \leftarrow \mathbf{a} + \mathbf{y}_k$ 
4: until all patterns properly classified
5: return  $a$ 
```

Algorithm 5.8 (Batch Relaxation with Margin)

```
1: initialize  $a, \eta(\cdot), b, k \leftarrow 0$ 
2: do  $k \leftarrow (k + 1) \bmod n$ 
3:    $\mathcal{Y}_k = \{\}$ 
4:    $j = 0$ 
5:   do  $j \leftarrow j + 1$ 
6:     if  $\mathbf{a}^t \mathbf{y}^j \leq b$  then Append  $\mathbf{y}^j$  to  $\mathcal{Y}_k$ 
7:   until  $j = n$ 
8:    $\mathbf{a} \leftarrow \mathbf{a} + \eta(k) \sum_{\mathbf{y} \in \mathcal{Y}_k} \frac{b - \mathbf{a}^t \mathbf{y}}{\|\mathbf{y}\|^2} \mathbf{y}$ 
9: until  $\mathcal{Y}_k = \{\}$ 
10: return  $a$ 
```

Results

Long training times proved problematic, due to the fact that it took greater than 100000 iterations to reach convergence using the batch relaxation rule for both data sets. Most likely, this is due to the fact that the mean and standard deviation of the training samples was not normalized (however, the perceptron was trained using normalized augmented features, e.g. for training ω_1 , $y_i^T = [1, \mathbf{x}_i]$ if x_i belongs to ω_1 , and $y_i^T = [-1, -\mathbf{x}_i]$ if x_i is not a sample for ω_1 .

In order to compensate for the long training times, training would terminate after 100000 trials or if the change in weights was approximately zero and $J_p(a) \approx 0$ (Equation (33) from Chapter 5). It's very likely that these early-termination heuristics had a detrimental effect on the classification accuracy, since in these cases, there was no guarantee that $\mathbf{a}^T \mathbf{y} \geq 0$ for all augmented \mathbf{y} .

UCI Wine Data Set

Algorithm 5.4 (Fixed-Increment Single-Sample Perceptron)

Training Statistics		
Class	Iterations	Runtime (s)
ω_1	22864	27.498
ω_2	52918	62.130
ω_3	35600	127.227
Total	111,382	216.855

The weights for each class after training were:

$$\begin{aligned}
 \mathbf{a}_1 &= \begin{bmatrix} 61035.869 \\ -29616.504 \\ 21752.882 \\ 1483.633 \\ -25062.26 \\ 157.921 \\ 11769.882 \\ 39925.061 \\ -8012.231 \\ 32849.055 \\ 18856.758 \\ -8917.463 \\ 31289.131 \\ 568.08 \end{bmatrix} & \mathbf{a}_2 &= \begin{bmatrix} 137199.142 \\ 37238.975 \\ -18303.661 \\ 6240.703 \\ 5099.608 \\ -3584.203 \\ 28903.859 \\ -3043.506 \\ 21158.897 \\ -25565.128 \\ -60256.757 \\ 34175.689 \\ 20013.206 \\ -55.379 \end{bmatrix} & \mathbf{a}_3 &= \begin{bmatrix} 90871.97 \\ 2481.881 \\ -12760.566 \\ -7305.273 \\ 6178.2 \\ 3980.622 \\ -31851.585 \\ -51828.543 \\ -3346.438 \\ -17933.398 \\ 39059.302 \\ -20064.289 \\ -62224.546 \\ -595.573 \end{bmatrix}
 \end{aligned}$$

This resulted in 78 correct classifications out of 89 (87.6% accuracy).

Algorithm 5.8 (Batch Relaxation with Margin)

Due to the large number of training iterations, training was capped to 100000 iterations.

Training Statistics

Class	Iterations	Runtime (s)
ω_1	100000	104.886
ω_2	100000	146.686
ω_3	100000	147.255
Total	300000	398.827

After training the perceptron using Algorithm 5.8, the weights were:

$$\mathbf{a}_1 = \begin{bmatrix} 0.582 \\ -0.213 \\ 0.487 \\ 0.728 \\ -0.381 \\ -0.069 \\ 0.881 \\ 0.103 \\ 0.246 \\ 0.188 \\ 0.494 \\ 0.443 \\ 0.534 \\ 0.008 \end{bmatrix} \quad \mathbf{a}_2 = \begin{bmatrix} 0.539 \\ 0.306 \\ 0.402 \\ 0.633 \\ -0.193 \\ -0.021 \\ 0.284 \\ 0.145 \\ 0.69 \\ 0.973 \\ -0.288 \\ 0.054 \\ 0.657 \\ -0.006 \end{bmatrix} \quad \mathbf{a}_3 = \begin{bmatrix} 0.371 \\ 0.141 \\ 0.511 \\ 0.841 \\ -0.176 \\ -0.042 \\ 0.062 \\ 0.114 \\ 0.545 \\ 0.337 \\ 0.601 \\ 0.142 \\ -0.14 \\ -0.003 \end{bmatrix}$$

This resulted in 83 correct classifications out of 89 (93.26% accuracy)

USPS Handwritten Digit Data Set

Algorithm 5.4 (Fixed-Increment Single-Sample Perceptron)

Training Statistics

Class	Iterations	Runtime (s)
ω_0	15	0.061
ω_1	5	0.021
ω_2	13	0.049
ω_3	9	0.033
ω_4	24	0.081
ω_5	17	0.050
ω_6	11	0.033
ω_7	21	0.062
ω_8	20	0.059
ω_9	42	0.118
Total	177	0.571

The weights are too large to display in the report, however they can be displayed by using the run.py script, by invoking the command “python3 run.py fixed bin/digits_train.txt bin/digits_test.txt”

This resulted in 1598 correct classifications out of 2007 (79.62% accuracy)

Algorithm 5.8 (Batch Relaxation with Margin)

In order to reach convergence more quickly, the algorithm terminated when the change in weights was approximately zero, and $J_r(a) \approx 0$ (Equation (33) from Chapter 5).

Training Statistics

Class	Iterations	Runtime (s)
ω_0	1114	2.923
ω_1	1424	3.833
ω_2	1106	2.955
ω_3	1179	3.351
ω_4	1357	3.856
ω_5	1524	4.176
ω_6	1331	3.617
ω_7	2335	6.104
ω_8	1189	3.186
ω_9	2561	7.022
Total	15120	41.023

The weights are too large to display in the report, however they can be displayed by using the run.py script, by invoking the command “python3 run.py relax bin/digits_train.txt bin/digits_test.txt”

This resulted in 1614 correct classifications out of 2007 (80.42% accuracy)

2 Multi-Class Classification

Use the basic two-class perceptron algorithms to solve multi-class classification problems by using the one-against-the-rest and one-against-the-other methods. Note that you need to handle ambiguous cases properly.

Results

TODO For each dataset, now train a classifier to classify all the classes using the one-against-the-rest and the one-against-the-other methods based on the two two-class algorithms, resulting in four different classifiers on each dataset and then classify the test set. Document classification accuracy, iterations in training, and classification time for test, and compare the one-against-the-rest and the one-against-the-other methods.

UCI Wine Data Set

USPS Handwritten Digits Data Set

3 Adaboost to Create Strong Classifiers

Implement Algorithm 1 (AdaBoost) in Chapter 9 of the textbook to create a strong classifier using the above linear discriminant functions.

Algorithm 9.1 (AdaBoost)

```
1: initialize  $\mathcal{D} = \{\mathbf{x}^1, y_1, \dots, \mathbf{x}^n, y_n\}, k_{max}, W_1(i) = 1/n, i = 1 \dots n$ 
2:  $k = 0$ 
3: do  $k \leftarrow k + 1$ 
4:   train weak learner  $C_k$  using  $\mathcal{D}$  sampled according to  $W_k(i)$ 
5:    $E_k \leftarrow$  training error of  $C_k$  measured on  $\mathcal{D}$  using  $W_k(i)$ 
6:    $\alpha_k \leftarrow 0.5 \ln [(1 - E_k)/E_k]$ 
7:    $W_{k+1}(i) = \frac{W_k(i)}{Z_k} \times \begin{cases} e^{-\alpha_k} & \text{if } h_k(\mathbf{x}^i) = y_i \text{ (correct classification)} \\ e^{\alpha_k} & \text{if } h_k(\mathbf{x}^i) \neq y_i \text{ (incorrect classification)} \end{cases}$ 
8: until  $k = k_{max}$ 
9: return  $C_k$  and  $\alpha_k$  for  $k = 1$  to  $k_{max}$  (ensemble of classifiers with weights)
```

Results

Boost Algorithm 8 to create a strong classifier for class 1 vs. class 2, class 1 vs. class 3, and class 2 vs. class 3 on the two datasets. Then classify the corresponding test samples from the relevant classes in test sets (in other words, for example, for the class 1 vs. class 2 classifier, you only need to classify test samples from classes 1 and 2); then document classification accuracy and show and analyze the improvement.

UCI Wine Data Set

USPS Handwritten Digits Data Set

4 Extra Credit

4.1 Support vector machines

By using an available quadratic programming optimizer or an SVM library, implement a training and classification algorithm for support vector machines. Then use your algorithm on the USPS dataset. Document the classification accuracy and compare the results with that from the two basic algorithms.

Results

TODO

4.2 Kernel method for linear discriminant functions

Given a kernel function, derive the kernel-version of Algorithm 4 and implement the algorithm, and then apply it on the given wine and USPS datasets. Document the classification accuracy and compare the results with that from the two basic algorithms without kernels. Use the polynomial function of degree three as the kernel function; optionally, you can use other commonly used kernel functions.

Results

TODO

4.3 Multiple-class linear machines and multiple-class boosting

Use the Keslers construction to train a linear machine for multi-class classification and then use the SAMME algorithm to boost its performance on the training set. Apply the algorithm on both datasets and classify the corresponding test samples in the test sets. Document the classification accuracy and compare the results with that from the one-against-the-rest and one-against-the- other algorithms.

4.3.1 Results

TODO

References

- [1] Jones E, Oliphant E, Peterson P, *et al.* **SciPy: Open Source Scientific Tools for Python**, 2001-, <http://www.scipy.org/> [Online; accessed 2015-10-24].